

# node.js

ry@tinyclouds.org

April 17, 2010

# **node.js** : Evented Server-side Javascript

Good at handling lots of different kinds of I/O at the same time.

Achieves this by all making network I/O nonblocking and all file I/O asynchronous. (Almost all.)

```
1  http = require('http');
2  net = require('net');
3  c = 0;
4
5  http.createServer(function(req, res) {
6      c++;
7      res.writeHead(200);
8      res.end('hello world\n');
9  }).listen(8000);
10
11 net.createServer(function(socket) {
12     socket.write('connections: ' + c);
13     socket.end();
14 }).listen(8001);
```

# For a detailed introduction see

[http://jsconf.eu/2009/video\\_nodejs\\_by\\_ryan\\_dahl.html](http://jsconf.eu/2009/video_nodejs_by_ryan_dahl.html)

# Speed

# Benchmarks


- ▶ nginx v0.7.65
- ▶ node v0.1.91
- ▶ tornado v0.2 (python 2.6.4)
- ▶ thin v1.2.7 (ruby 1.9.1-p376)

In Linux using a Intel Core 2 Due  
2.53, 4 GB memory

The standard 'hello world'  
concurrency benchmark.


(Approximately 100 byte response  
for each.)

## Smaller is Better





## Smaller is Better




How do the servers perform when the concurrency is fixed at 300, but serve different response sizes.


This is what the node version looks like, approximately

```
1  string = '';  
2  for (i = 0; i < 16*1024; i++) {  
3      string += 'd';  
4  }  
5  
6  http.createServer(function(req, res) {  
7      res.writeHead(200);  
8      res.end(string, 'ascii');  
9  });
```

concurrency=300, Smaller is Better



concurrency=300, Smaller is Better



Wow. Node sucks at serving large files.

Well over 3 second responses for 256 kilobyte files at 300 concurrent connections.

What's happening:

V8 has a generational garbage collector. Moves objects around randomly.


Node can't get a pointer to raw string data to write to socket.

Using Node's new **Buffer** object,  
the results change.



```
1  buffer = new Buffer(16*1024);
2  for (i = 0; i < buffer.length; i++) {
3      buffer[i] = 100;
4  }
5
6  http.createServer(function(req, res) {
7      res.writeHead(200);
8      res.end(buffer);
9  });
```


concurrency=300, Smaller is Better



Node can push a **Buffer** reference to socket pretty quickly.

A histogram of the 256 kilobyte case:

c=300, res=256kb



But the fact remains, pushing large strings to socket is slow.

Hopefully this can be mitigated in the future.

```
// Create
new Buffer(size);

// Modify
buffer[i] = value;

// memcpy
bufferA.copy(bufferB, offset);

// encode a string
buffer.write('string', 'utf8', offset);

// decode a string
buffer.toString('utf8', offset);
```

Since the 0.1.90 release, most of Node is written in JavaScript.

Just thin bindings.

Part of this rewrite was to expose the raw memory **Buffer** to users.

The other part was to unify Streams.



Node had all of these objects which would stream data, but had different interfaces.

HTTP request object: emitted a **'body'** event for each new chunk of data, and had a **sendBody()** method.

TCP socket: emitted a '**receive**' event for each new chunk of data, and had a **send()** method.

Stdio: emitted a '**data**' event for each chunk on stdin, and had a **write()** method for stdout.

**'data' , 'receive' , 'body'**

**sendBody() , send() ,  
write()**

I slowly realized that I should stop making up different names for these things.

With a unified interface a general purpose ‘pumping’ function, with all the proper **throttling** and error checking is possible!

```
pump(request, stdout, callback);
```

```
pump(request, file, callback);
```

```
pump(file, response, callback);
```

Stream interface is split into two parts: Readable and writable streams.

Some streams are both readable and writable.

# Readable stream.

- ▶ Event: **'error'**
- ▶ Event: **'data'**
- ▶ Event: **'end'** (EOF or FIN)
- ▶ **pause()**
- ▶ **resume()**
- ▶ **destroy()**



# Writable stream.

- ▶ Event: **'error'**
- ▶ Event: **'drain'**
- ▶ **write()**
- ▶ **end()**
- ▶ **destroy()**

Readable:

- ▶ `stdin`
- ▶ `ServerRequest`
- ▶ `childProcess.stdout`

Writable:

- ▶ `stdout`
- ▶ `ServerResponse`
- ▶ `childProcess.stdin`

Probably need to expand the interface for more complex types of throttling.

(E.G. a low water mark for when to draining.)

# Other Improvements.

For technical reasons:

- ▶ UDNS has been replaced with C-Ares.
- ▶ GnuTLS has been replaced with OpenSSL.

Side effect: Node no longer depends on any GPL or LGPL libraries.

The REPL library supports connecting to arbitrary sockets, not just stdio.

So you can setup a UNIX socket server at

**/tmp/my\_node\_app.sock** and 'telnet' into it with the **socat** utility.

```
1 net.createServer(function (socket) {  
2   repl.start('my app> ', socket);  
3 }).listen('/tmp/my_node_app.sock');
```

**promise.wait()** was removed.

This was a kind of cooperative  
threading thing. Execution would  
jump to a new stack when called.

Coroutines complicate the mental model while adding only cheap syntactic pleasures.

Must worry about I/O occurring in all function calls. (They might call **wait()**.) The user needs to make their functions coroutine safe!



Cooperative threading of any sort is  
a bad idea.


# Build Bot



# Build Bot


## Linux AMD64

### Static Http Server




Legend: t  
move mouse over graph Shift-click to place baseline

### Timers



Legend: t  
rb98cd6753b4: 2827.00 ms +/- Shift-click to place


### Process Loop



Legend: t  
move mouse over graph Shift-click to place baseline


## Linux Xeon

### Static Http Server




Legend: t  
move mouse over graph Shift-click to place baseline

### Timers



Legend: t  
move mouse over graph Shift-click to place baseline

### Process Loop



Legend: t  
move mouse over graph Shift-click to place baseline

# The Project Is Growing

42 releases. Current version: v0.1.91

63 contributors.

6000 lines of JavaScript, 11000 lines  
of C++

1100 people on the mailing list, 1400  
watchers on GitHub.

`http://nodejs.org/`

`http://wiki.github.com/ry/  
node/modules`

`http://howtonode.org/`

**`ry@tinyclouds.org`**